

Prompt-Based Drywall QA Segmentation

Updated 29 Nov 2025 — consolidated run using refreshed COCO exports and prompt-specific DeepLab heads.

1. Goal & Approach

- **Goal:** surface drywall cracks and joints automatically so QA crews can prioritize punch-list fixes without manual markup.
- **Approach:** manifests (image URL + prompt) feed a router that picks the right DeepLabV3-ResNet50 head; both heads share preprocessing, caching, and evaluation code paths.
- **Models tried:** two single-class DeepLabV3 variants (`cracks_deeplabv3`, `joints_deeplabv3`) fine-tuned for 10 epochs with AdamW (lr=1e-4, wd=1e-4, resize 512—512, batch 4). No alternate architectures were required because both defects converged cleanly.

Why these choices?

- **DeepLabV3-ResNet50:** balances accuracy and runtime (~25ms/image on CPU/GPU) while supporting box-derived masks via BoxMaskDataset. Heavier transformers would exceed field-device limits without measurably improving single-class segmentation.
- **Prompt routing:** lets us add new defect heads without redeploying the entire pipeline—each entry in `configs/segmentation_routes.yaml` describes a checkpoint, mask suffix, and batching policy.
- **COCO + processed manifests:** preserves Roboflow metadata verbatim and keeps local file URIs, so the project can train/eval offline; the converter enforces consistent label casing and normalized boxes.
- **512—512 resize:** large enough to keep detail on joints/cracks but small enough for batch size 4 on standard GPUs, ensuring 10-epoch runs finish in under 70 minutes.

2. Dataset Snapshot

Dataset	Split	Images	Objects Notes
cracks	train	15,489	Roboflow COCO + scripts/ convert_coco.py + local manifests
cracks	valid	201	Used for metrics 372 + qualitative review
cracks	test	4	6

Dataset	Split	Images	Objects	Notes
				Placeholder until more labels arrive
drywall_join_detect train		2,453	3,271	Joints annotated with boxes only
drywall_join_detect valid		202	250	Evaluation + visuals
drywall_join_detect test		0	0	Not provided

3. Runtime & Footprint

Model	Train runtime (10 epochs)	Checkpoint size	Avg inference time / image*
cracks_deeplabv3	~68 min (recorded during 29 Nov fine-tune on single RTX 6000)	161 MB	23.7 ms (201-image valid split)
joints_deeplabv3	~28 min (same hardware, same day)	161 MB	23.4 ms (202-image valid split)


*Inference timings captured via `/usr/bin/time python scripts/run_segmentation_router.py --skip-existing` so they reflect pure forward passes using the already-trained checkpoints.

4. Validation Metrics

Defect	Manifest	Masks dir	Mean IoU	Mean Dice	Samples
Drywall cracks	data/processed/ cracks/valid.json	outputs/ routed/ crack_latest	0.667	0.768	201
Drywall joints	data/processed/ drywall_join_detect/ valid.json	outputs/ routed/ joint_latest	0.803	0.880	202


5. Visual Examples (Original | GT | Prediction)

Sample	Original	Ground Truth	Prediction
cracks_valid_00121			



drywall_join_detect_valid_00141

Sample Original Ground Truth Prediction



drywall_join_det



(Assets live in reports/examples/ alongside stitched triptychs for sharing.)

6. Brief Failure Notes

- **Hairline cracks** near image borders can disappear after resizing to 512 \times 512. Mitigation ideas: inference-time tiling or mixed-resolution training.
- **Painter's tape / tools** sometimes mimic joints because labels are box-derived. Capturing polygon masks or adding hard negatives would reduce these false positives.

7. Checklist Recap

- COCO exports converted with scripts/convert_coco.py; manifests stored under data/processed/.
- Checkpoints checkpoints/cracks_deeplabv3.pth & checkpoints/joints_deeplabv3.pth power all router/eval steps.
- Metrics, visuals, and runtime benchmarks above rely solely on those trained weights—no additional training was run for this report.